



# Vectorizing common HEP analysis algorithms

Jaydeep Nandi

Jim Pivarski

David Lange

# Contents

1. Definitions
2. Vectorization
3. Autovectorization
4. Add as the sections are defined.



3

# Definitions

# Definitions

- ▀ Parents:
- ▀ Offsets:
- ▀ Starts:
- ▀ Events:
- ▀ Stops:

# Scalar Programming and Vectorization

# Scalar Programming

- Works on scalar ( individual elements) values of an array.
- Applies an operation over the elements via a loop, usually a **for-loop** or **while** loop.
- Low-level, and tedious to write.
- Virtually all programming languages support scalar programming.

# Scalar Programming

- **For example:** If we wish to add two compatible vectors ( of same shape and castable type) **A** and **B** together and store it in **C**, a typical scalar program (in Python) would look like:

```
for i in range(len(A)):  
    C[i] = A[i] + B[i]
```

- It operates on each element of the arrays sequentially.

# What is Vectorization?

- Converting a scalar program to a vector program.
- It is also known as Array Programming.
- Wikipedia Definition: “In computer science, **array programming languages** (also known as vector or **multidimensional** languages) generalize operations on scalars to apply transparently to vectors, matrices, and higher-dimensional arrays.”
- Allows applying an operation on multiple data items simultaneously



# Array Languages

- Some languages and libraries support vectorization by default. Examples include:
  - **Python: Numpy, Scipy etc.**
  - **MATLAB**
  - **GNU OCTAVE**
  - **R etc.**
- Usually they compute the vectors under the hood as efficient “**C**” or “**FORTRAN**” implementations.

# Array Languages

- ▶ Recollect the addition of all elements of two vectors to give a new vector, which we implemented as a scalar code:

```
for i in range(len(A)):  
    C[i] = A[i] + B[i]
```

- ▶ In an array language, it is quite simple to write:

```
C = A + B
```

# So why Vectorize?

- Simple to write; no need of writing loops incessantly.
- Expressive, from a mathematical point of view.
- But those are not the major reason for vectorising code.
  - Most modern CPUs and GPUs provide support for vectorised code, which runs very efficiently, operating in a SIMD style.
  - Can take advantage of SSE ( Streaming SIMD Extensions) and AVX instructions, which operate on 4, 8 or more data simultaneously.
  - GPUs take it even further: Can operate on a large number, typically in thousands, of data at once. ( More on it later!)

# Vectorization and GPUs

# General Purpose GPU (GPGPU) Programming

- Works on Single Instruction Multiple Thread (SIMT) architecture.
- Processes multiple (depending on the number of threads available) data elements simultaneously.
- However, it requires parallel instructions, which don't have a dependency that can create a data race.
- Can be programmed with CUDA for Nvidia cards, and OpenCL for AMD cards and others.
- Primarily in C; bindings exist for other languages like PyCUDA and PyOPENCL for Python.

# Vectorization and GPGPU programming

- Vectorized codes are inherently parallel. They apply same instructions on multiple data elements of the vector at once.
- Is in perfect agreement with GPGPU philosophy.
- Vectorized codes thus give very high speedups if operated on the GPU.

# Autovectorization

# Autovectorization

- Modern compilers allow the vectorization of simple sequential loops into efficient SIMD instructions for CPU, typically through a compiler switch.
- Some compilers which support this:
  - **GCC**: Supports through switch **-free-vectorize** or **-O3** level optimization.
  - **ICC**: From Intel. Reported to be better than GCC at vectorising code.
  - **MSVC and others**.
- They check if a loop is safe to be vectorised, and if it is, they vectorise the code. The list of such possible vectorizable loops for GCC can be found [here](#).



# Autovectorization

- As an example, consider again the addition of two vectors. A simple c loop that does is :

```
for ( int i=0; i<length(A); i++)  
    C[i] = A[i] + B[i];
```

- The corresponding assembly dump ( relevant part) is:

```
vmovdqu xmm0, XMMWORD PTR [r9+rax]  
add edx, 1  
vinserti128 ymm0, ymm0, XMMWORD PTR [r9+16+rax], 0x1  
vpadd ymm0, ymm0, YMMWORD PTR [r13+0+rax]  
vmovups XMMWORD PTR [rcx+rax], xmm0
```

- Note the vector instructions in the assembly ( **vmovdqu**, **vinserti** etc)

# Autovectorization

- Advantages:
  - Sufficiently simpler than writing SSE or AVX instructions.
  - Easy to debug.
  - Portable.
- There is however, a major disadvantage. They cannot vectorise a loop if there is a loop carried dependency.

# Barriers to Autovectorization

- The major barrier is a loop carried dependency. They are variables whose particular value depends on the order of the execution.
- As an example consider the problem of finding the max value of all elements in an array **A**.

```
for (int i=0; i<N; i++)  
    maxval = max(maxval, A[i]);
```

- The value in maxval at any instance depends on the order of loop execution. So, it's a loop carried dependency, and the compiler cannot autovectorize it.

# Barriers to Autovectorization

- The assembly output of the loop is

```
mov edx, DWORD PTR [rax]
add rax, 4
cmp ecx, edx
jl .L12
cmp rax, rsi
jne .L11
mov edx, ecx
```

- Note the absence of vector instructions in the assembly, indicating an unvectorized loop.
- The reduction operation can be vectorised, and will be dealt with shortly.

# Vectorization of HEP analysis algorithms

# Algorithms Considered

- Argproduct ( argcombinations)
  - Local Reduction
  - And others.
- 
- For the complete list of algorithms and their detailed analysis, please refer to the complete report available [here](#).
  - Note that the algorithms are not physics-specific algorithms per se, but primitives which can be used to build them. They are meant to serve the purpose of demonstrating the efficiency of vectorization.

# Argproduct

# Argproduct

- A typical application in HEP is operating on the combinations or pairs of different particles in an event. They are useful for calculating different metrics like the deltaR distance, and in generated-reconstructed particle matching.
- Argproduct returns two arrays, consisting of the index of first element in the pair, and another array consisting of the index of the second element in the pair.
- Argproduct isn't inherently vectorizable, it has a loop carried dependency.



# Argproduct

- Sequential python code:

```
for i in range(events):  
    pairs_i = 0  
    for j in range(starts1[i], stops1[i]):  
        for k in range(strats2[i], stops2[i]):  
            first[pairs_i] = j  
            second[pairs_i] = k
```

- Note that `pairs_i` is a loop carried dependency.

# Argproduct

- ▶ The code can be vectorised if we have the information of
  - ▶ **Parents** array of every pair.
  - ▶ The number of elements in either array, say given by **counts** array, and the **starts** and **stops** arrays for the two particle sets.
  - ▶ A running index array for all the pairs in all the events.
- ▶ An important thing to notice is that if we consider the linear indexed arrays as matrices, then the first and second arrays are given by the row index and column index of the matrix.
- ▶ So, we just need a way to convert the linear index to matrix index.

# Argproduct

- For a simple illustration, consider the two arrays

```
arr1 = ['a', 'b', 'c']  
arr2 = [1,2]
```

- The output of argproduct will be

```
first = [0,0,1,1,2,2]  
second = [0,1,0,1,0,1]
```

- Which correspond to the pairs

```
[['a', 1], ['a', 2], ['b', 1], ['b', 2], ['c', 1], ['c', 2]]
```

# Argproduct

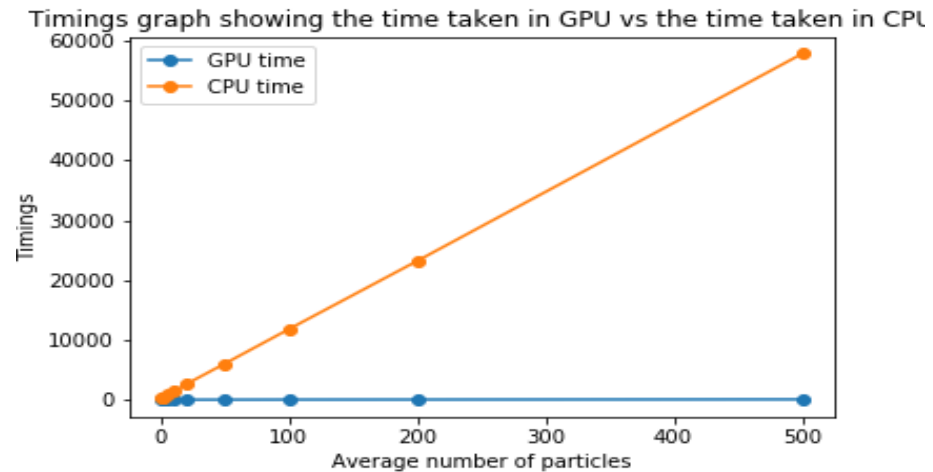
- Now, there are  $3 \times 2 = 6$  possible pairs. Notice that if we have a running index array:  
`index = [0,1,2,3,4,5]`
- Then **first** can be generated as the integer division of elements in **index** with number of elements in second array, given by **counts**, and **second** is the result of **modulo** operation on index by **counts**.

```
first = index // counts  
second = index % counts
```

- The operations are vectorised.
- The per-event case then can be derived by giving the proper offset to the indices.

# Performance Improvements of argproduct

- Vectorizing the argproduct gives a high speedup, especially when considered between GPU and CPU.



# Local Reduction

# Local Reduction

- Reduces an array per-event. The reduction operator can be any associative operator like **max()**, **min()** or **sum()**.
- Significantly tough to vectorise.
- We shall consider `max()`, and the rest can be simply substituted in place of `max()`
- Sequential code

```
for i in range(events):  
    for i in range(starts[i], stops[i]):  
        val = max(val, A[i])
```

# Local reduction

- The reduction has been implemented in parallel through a modified version of Hillis-Steele algorithm.
- Let us consider a single event first. The pseudocode for the upsweep phase of the algorithm is

```
for d=0 to ( $\log_2 n - 1$ ) do  
  forall k=0 to n-1 by  $2^{d+1}$  do  
     $a[k+2^{d+1}-1] = \max(a[k+2^d-1], a[k+2^{d+1}-1])$ 
```

- The algorithm proceeds by tree reducing the array elements, until a single element remains.



# Local Reduction

- To extend it to per-event case, we will again require the parents array for the data.
- It will serve as a mask for

# Vectorized reduction

# Vectorized local reduction

- A classic problem, that can be solved with a variation of Hillis-Steele algorithm.

- Todo next:
  - Vectorization and GPUs
  - Examples
    - Product
    - Local Reduction
  - Link to markdown report
  - Conclude